

```

!sudo apt update
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://dlcdn.apache.org/spark/spark-3.3.1/spark-3.3.1-bin-hadoop3.tgz

Get:1 https://cloud.r-project.org/bin/linux/ubuntu focal-cran40/ InRelease [3,622 B]
Get:2 https://cloud.r-project.org/bin/linux/ubuntu focal-cran40/ Packages [71.1 kB]
Ign:3 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 InRelease
Hit:4 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2004/x86_64 InRelease
Hit:5 http://archive.ubuntu.com/ubuntu focal InRelease
Get:6 http://security.ubuntu.com/ubuntu focal-security InRelease [114 kB]
Get:7 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal InRelease [18.1 kB]
Hit:8 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 Release
Get:9 http://archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
Hit:11 http://ppa.launchpad.net/cran/libgit2/ubuntu focal InRelease
Hit:12 http://ppa.launchpad.net/deadsnakes/ppa/ubuntu focal InRelease
Get:13 http://archive.ubuntu.com/ubuntu focal-backports InRelease [108 kB]
Hit:14 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu focal InRelease
Get:15 http://archive.ubuntu.com/ubuntu focal-updates/universe amd64 Packages [1,291 kB]
Get:16 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal/main Sources [2,381 kB]
Get:17 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 Packages [2,921 kB]
Get:18 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal/main amd64 Packages [1,128 kB]
Fetched 8,150 kB in 6s (1,280 kB/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
29 packages can be upgraded. Run 'apt list --upgradable' to see them.

!tar xf spark-3.3.1-bin-hadoop3.tgz
!pip install -q findspark
!pip install pyspark

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting pyspark
  Downloading pyspark-3.3.1.tar.gz (281.4 MB)
    281.4/281.4 MB 4.7 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Collecting py4j==0.10.9.5
  Downloading py4j-0.10.9.5-py2.py3-none-any.whl (199 kB)
    199.7/199.7 KB 23.0 MB/s eta 0:00:00
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.3.1-py2.py3-none-any.whl size=281845512 sha256=a6afb6da4b93660e4f77288dbc469dc149ff
  Stored in directory: /root/.cache/pip/wheels/43/dc/11/ec201cd671da62fa9c5cc77078235e40722170ceba231d7598
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9.5 pyspark-3.3.1

```

---

```

import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.3.1-bin-hadoop3"

from google.colab import files
# upload the Vector Assembler model av_model.zip
files.upload()
# upload the ML model rf_model.zip
files.upload()

# extract the models
import os
import zipfile
dirname = os.getcwd()
print('path=',dirname)

rf_filename = "rf_model.zip"
va_filename = "va_model.zip"

rf_path= os.path.join(dirname, rf_filename)
va_path = os.path.join(dirname, va_filename)

zip_rf = zipfile.ZipFile(rf_path, "r")
zip_rf.extractall(dirname)
zip_va = zipfile.ZipFile(va_path, "r")
zip_va.extractall(dirname)

zip_rf.close()
zip_va.close()

path= /content

```

```
import findspark
findspark.init()
findspark.find()
```

```
'/content/spark-3.3.1-bin-hadoop3'
```

```
from pyspark.sql import DataFrame, SparkSession
from typing import List
import pyspark.sql.types as T
import pyspark.sql.functions as F
from pyspark.sql.functions import isnull, when, count, col
from pyspark import SparkFiles
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
```

```
spark = SparkSession.builder \
    .master('local') \
    .appName("NY Parking Violation") \
    .config("spark.sql.adaptive.enabled","true") \
    .config("spark.executor.memory","10g") \
    .config("spark.driver.memory","10g") \
    .getOrCreate()

spark
```

#### SparkSession - in-memory

#### SparkContext

[Spark UI](#)

Version

v3.3.1

Master

local

AppName

NY Parking Violation

```
columns_selected = ["Registration State","Plate Type",\
    "Violation Code", "Vehicle Body Type","Vehicle Make","Issuing Agency", "Street Code1", \
    "Street Code2","Street Code3","Violation Location","Violation Precinct", \
    "Issuer Precinct","Issuer Code","Issuer Command",\
    "Violation County","Law Section","Sub Division","Vehicle Color"]
```

```
# cols_index = [df.toPandas().columns.get_loc(col) for col in columns_selected]
# print(cols_index)
```

```
from pyspark.sql import SparkSession
from pyspark.ml import PipelineModel
from pyspark.ml.feature import VectorAssembler, VectorIndexer, OneHotEncoder, StringIndexer
def preprocess_data(df):
    df.show(5)
    df.count(), len(df.columns)
    df = df.select(columns_selected)

    # clean up the data as many have incorrect values.
    df = df[(df['Registration State'] != "99") \
        & (df['Plate Type'] != "999") \
        & (df['Violation Code'] != 0)]
    # clean up the data
    # Check if the null value still exist
    df.select([count(when(isnull(c), c)).alias(c) for c in df.columns]).show()

    df = df.na.drop()
    df.dropDuplicates()

    # convert to required type
    cols = [F.col(field[0]).cast('double') if (field[1] == 'int') else F.col(field[0]) for field in df.dtypes]
    df = df.select(cols)

    #use model to transform
    pm = PipelineModel.load("/content/va_model")

    df = pm.transform(df)
    return df

def predict_data(df):
    #use ml model to predict
    pred_model = PipelineModel.load("/content/rf_model")
```

```
data = df.select(F.col("features_scaled").alias("features"))
# use the PipelineModel object to perform prediciton on data.
prediction = pred_model.transform(data)

# print the results
# prediction.select('label','prediction','Violation_Location').show(5)
prediction.select('prediction','Violation_Location').show(10)
return prediction
```

```
df = spark.read.csv("/content/PV.00002.csv",inferSchema=True, header= True)
df.count(),len(df.columns)
```

(100000, 51)

```
vec_df = prepreocess_data(df)
# predict_data(vec_df)
```

Summons Number	Plate ID	Registration State	Plate Type	Issue Date	Violation Code	Vehicle Body Type	Vehicle Make	Issuing Agency	Stre
7021662099	GJH8997	NY	PAS	08/27/2013	37	SUBN	SUBAR	T	
7021662105	GJU4876	NY	PAS	08/27/2013	37	SUBN	ACURA	T	
7022231175	5F15B	NY	OMT	07/29/2013	21	TAXI	FORD	T	
7022235820	GAC6868	NY	PAS	08/20/2013	20	SUBN	HONDA	T	
7022235879	88696JU	NY	COM	08/20/2013	14	VAN	CHEVR	T	

only showing top 5 rows

Registration State	Plate Type	Violation Code	Vehicle Body Type	Vehicle Make	Issuing Agency	Street Code1	Street Code2	Street Code3
0	0	0	52	51	0	0	0	0

```
vec_df.show(5,False)
```

Registration State	Plate Type	Violation Code	Vehicle Body Type	Vehicle Make	Issuing Agency	Street Code1	Street Code2	Street Code3
NY	PAS	37.0	SUBN	SUBAR	T	40690.0	24090.0	24140.0
NY	PAS	37.0	SUBN	ACURA	T	40690.0	24090.0	24140.0
NY	OMT	21.0	TAXI	FORD	T	23290.0	17790.0	18040.0
NY	PAS	20.0	SUBN	HONDA	T	65590.0	28990.0	52090.0
NY	COM	14.0	VAN	CHEVR	T	59990.0	16540.0	16790.0

only showing top 5 rows

```
prediction = predict_data(vec_df)
```

prediction	Violation_Location
4.0	109.0
4.0	109.0
35.0	103.0
35.0	103.0
27.0	104.0
35.0	103.0
27.0	104.0
25.0	108.0
35.0	103.0
27.0	104.0

only showing top 10 rows

```
prediction.show(10,True)
```

features	rawPrediction	probability	prediction	Violation_Location
(1926,[0,59,100,3...]	[0.55610398108441...]	[0.01390259952711...]	4.0	109.0
(1926,[0,59,100,3...]	[0.63189684659933...]	[0.01579742116498...]	4.0	109.0
(1926,[0,62,109,3...]	[0.43424074714016...]	[0.01085601867850...]	35.0	103.0
(1926,[0,59,100,3...]	[0.39215854643541...]	[0.00980396366088...]	35.0	103.0
(1926,[0,60,99,30...]	[0.30858237046250...]	[0.00771455926156...]	27.0	104.0
(1926,[0,63,100,3...]	[0.29918287285738...]	[0.00747957182143...]	35.0	103.0
(1926,[12,59,102,...]	[0.18392480567470...]	[0.00459812014186...]	27.0	104.0
(1926,[0,60,101,3...]	[0.01163596065937...]	[2.90899016484277...]	25.0	108.0
(1926,[6,59,100,3...]	[0.70451213965000...]	[0.01761280349125...]	35.0	103.0
(1926,[12,59,102,...]	[0.18392480567470...]	[0.00459812014186...]	27.0	104.0

+-----+  
only showing top 10 rows

```
df.show(10,True)
```

Summons Number	Plate ID	Registration State	Plate Type	Issue Date	Violation Code	Vehicle Body Type	Vehicle Make	Issuing Agency	Stre
7021662099	GJH8997	NY	PAS	08/27/2013	37	SUBN	SUBAR		T
7021662105	GJU4876	NY	PAS	08/27/2013	37	SUBN	ACURA		T
7022231175	5F15B	NY	OMT	07/29/2013	21	TAXI	FORD		T
7022235820	GAC6868	NY	PAS	08/20/2013	20	SUBN	HONDA		T
7022235879	88696JU	NY	COM	08/20/2013	14	VAN	CHEVR		T
7022235892	GAH6321	NY	OMS	08/20/2013	38	SUBN	TOYOT		T
7022235909	PAF8691	GA	PAS	08/21/2013	21	4DSD	CHEVR		T
7022235910	51976MA	NY	COM	08/21/2013	19	DELV	INTER		T
7022235934	928CA4	MA	PAS	08/21/2013	21	SUBN	ACURA		T
7022235983	PAF8691	GA	PAS	08/21/2013	38	4DSD	CHEVR		T

only showing top 10 rows

