```
!sudo apt update
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://dlcdn.apache.org/spark/spark-3.3.1/spark-3.3.1-bin-hadoop3.tgz
      Get:1 <a href="https://cloud.r-project.org/bin/linux/ubuntu">https://cloud.r-project.org/bin/linux/ubuntu</a> focal-cran40/ InRelease [3,622 B]
      Ign: 2 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 InRelease
     Hit:3 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2004/x86_64 InRelease
     Hit:4 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64
                                                                                                                  Release
     Get:5 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal InRelease [18.1 kB]
     Get:6 http://security.ubuntu.com/ubuntu focal-security InRelease [114 kB]
     Hit:7 http://archive.ubuntu.com/ubuntu focal InRelease
      Get:8 http://archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
     Hit:9 <a href="http://ppa.launchpad.net/cran/libgit2/ubuntu">http://ppa.launchpad.net/cran/libgit2/ubuntu</a> focal InRelease
      Get:10 https://cloud.r-project.org/bin/linux/ubuntu focal-cran40/ Packages [71.1 kB]
      Hit:11 <a href="http://ppa.launchpad.net/deadsnakes/ppa/ubuntu">http://ppa.launchpad.net/deadsnakes/ppa/ubuntu</a> focal InRelease
      Get:12 http://archive.ubuntu.com/ubuntu focal-backports InRelease [108 kB]
     Hit:14 <a href="http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu">http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu</a> focal InRelease
     Get:15 <a href="http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu">http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu</a> focal/main Sources [2,381 kB]
     Get:16 <a href="http://archive.ubuntu.com/ubuntu">http://archive.ubuntu.com/ubuntu</a> focal-updates/universe amd64 Packages [1,291 kB]
     Get:17 <a href="http://archive.ubuntu.com/ubuntu">http://archive.ubuntu.com/ubuntu</a> focal-updates/main amd64 Packages [2,921 kB]
     Get:18 <a href="http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu">http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu</a> focal/main amd64 Packages [1,128 kB]
      Fetched 8,150 kB in 4s (2,270 kB/s)
      Reading package lists... Done
     Building dependency tree
      Reading state information... Done
      29 packages can be upgraded. Run 'apt list --upgradable' to see them.
!tar xf spark-3.3.1-bin-hadoop3.tgz
!pip install -q findspark
!pip install pyspark
      Looking in indexes: <a href="https://pypi.org/simple">https://us-python.pkg.dev/colab-wheels/public/simple/</a>
      Collecting pyspark
        Downloading pyspark-3.3.1.tar.gz (281.4 MB)
                                                       - 281.4/281.4 MB 4.6 MB/s eta 0:00:00
        Preparing metadata (setup.py) ... done
     Collecting py4j==0.10.9.5
        Downloading py4j-0.10.9.5-py2.py3-none-any.whl (199 kB)
                                                       - 199.7/199.7 KB 20.0 MB/s eta 0:00:00
      Building wheels for collected packages: pyspark
        Building wheel for pyspark (setup.py) ... done
        Created wheel for pyspark: filename=pyspark-3.3.1-py2.py3-none-any.whl size=281845512 sha256=76036277844fabfd3e340fdc62515232b0da
        Stored in directory: /root/.cache/pip/wheels/43/dc/11/ec201cd671da62fa9c5cc77078235e40722170ceba231d7598
      Successfully built pyspark
      Installing collected packages: py4j, pyspark
     Successfully installed py4j-0.10.9.5 pyspark-3.3.1
     4
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK HOME"] = "/content/spark-3.3.1-bin-hadoop3"
import findspark
findspark.init()
findspark.find()
      '/content/spark-3.3.1-bin-hadoop3'
from pyspark.sql import DataFrame, SparkSession
from typing import List
import pyspark.sql.types as T
import pyspark.sql.functions as F
from pyspark.sql.functions import isnull, when, count, col, isnan
from pyspark import SparkFiles
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
spark = SparkSession.builder \
         .master('local') \
         .appName("NY Parking Violation") \
         .config("spark.sql.adaptive.enabled","true") \
         .config("spark.executor.memory","10g") \
         .config("spark.driver.memory","10g") \
         .getOrCreate()
snark
```

```
SparkSession - in-memory
     SparkContext
     Spark UI
     Version
     v3.3.1
from google.colab import files
uploaded = files.upload()
     Choose Files No file chosen
                                         Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to
     enable
     Saving DV AAAA? cey to DV AAAA? cey
## ensure to use the link of the raw file of the csv
# url = "https://data.cityofnewyork.us/api/views/pvqr-7yc4/rows.csv"
# spark.sparkContext.addFile(url)
# df = spark.read.csv(SparkFiles.get("rows.csv"),inferSchema=True, header= True)
df = spark.read.csv("/content/PV.00001.csv",inferSchema=True, header=True)
df.count(),len(df.columns)
     (100000, 51)
df.printSchema()
     root
       |-- Summons Number: long (nullable = true)
       |-- Plate ID: string (nullable = true)
       |-- Registration State: string (nullable = true)
       -- Plate Type: string (nullable = true)
       -- Issue Date: string (nullable = true)
       |-- Violation Code: integer (nullable = true)
       -- Vehicle Body Type: string (nullable = true)
       -- Vehicle Make: string (nullable = true)
       -- Issuing Agency: string (nullable = true)
       |-- Street Code1: integer (nullable = true)
|-- Street Code2: integer (nullable = true)
       -- Street Code3: integer (nullable = true)
       -- Vehicle Expiration Date: integer (nullable = true)
       -- Violation Location: integer (nullable = true)
       |-- Violation Precinct: integer (nullable = true)
       -- Issuer Precinct: integer (nullable = true)
       |-- Issuer Code: integer (nullable = true)
       -- Issuer Command: string (nullable = true)
       |-- Issuer Squad: string (nullable = true)
       |-- Violation Time: string (nullable = true)
       -- Time First Observed: string (nullable = true)
       -- Violation County: string (nullable = true)
       -- Violation In Front Of Or Opposite: string (nullable = true)
       -- House Number: string (nullable = true)
       -- Street Name: string (nullable = true)
       -- Intersecting Street: string (nullable = true)
       -- Date First Observed: integer (nullable = true)
       -- Law Section: integer (nullable = true)
       -- Sub Division: string (nullable = true)
       -- Violation Legal Code: string (nullable = true)
       |-- Days Parking In Effect : string (nullable = true)
       -- From Hours In Effect: string (nullable = true)
       -- To Hours In Effect: string (nullable = true)
       -- Vehicle Color: string (nullable = true)
-- Unregistered Vehicle?: integer (nullable = true)
       -- Vehicle Year: integer (nullable = true)
       -- Meter Number: string (nullable = true)
       -- Feet From Curb: integer (nullable = true)
       |-- Violation Post Code: string (nullable = true)
       -- Violation Description: string (nullable = true)
       -- No Standing or Stopping Violation: string (nullable = true)
       |-- Hydrant Violation: string (nullable = true)
       -- Double Parking Violation: string (nullable = true)
       -- Latitude: string (nullable = true)
       |-- Longitude: string (nullable = true)
|-- Community Board: string (nullable = true)
       -- Community Council : string (nullable = true)
       -- Census Tract: string (nullable = true)
       -- BIN: string (nullable = true)
       -- BBL: string (nullable = true)
       -- NTA: string (nullable = true)
```

Analyze and preprocess the data

```
# Remove all columns which have more than 70% null values
null threshold = 0.70
columns valid = [x for x in df.columns if df.filter(F.col(x).isNull()).count()/df.count() < null threshold]</pre>
df = df.select(*columns_valid)
df.count(), len(df.columns)
    (100000, 37)
df.printSchema()
     root
      |-- Summons Number: long (nullable = true)
      |-- Plate ID: string (nullable = true)
      -- Registration State: string (nullable = true)
      |-- Plate Type: string (nullable = true)
      |-- Issue Date: string (nullable = true)
      |-- Violation Code: integer (nullable = true)
      -- Vehicle Body Type: string (nullable = true)
      -- Vehicle Make: string (nullable = true)
      -- Issuing Agency: string (nullable = true)
      -- Street Code1: integer (nullable = true)
      -- Street Code2: integer (nullable = true)
      |-- Street Code3: integer (nullable = true)
      |-- Vehicle Expiration Date: integer (nullable = true)
      |-- Violation Location: integer (nullable = true)
      |-- Violation Precinct: integer (nullable = true)
      |-- Issuer Precinct: integer (nullable = true)
      -- Issuer Code: integer (nullable = true)
      -- Issuer Command: string (nullable = true)
      -- Issuer Squad: string (nullable = true)
      |-- Violation Time: string (nullable = true)
      -- Violation County: string (nullable = true)
      -- Violation In Front Of Or Opposite: string (nullable = true)
      |-- House Number: string (nullable = true)
      -- Street Name: string (nullable = true)
      |-- Intersecting Street: string (nullable = true)
      -- Date First Observed: integer (nullable = true)
      |-- Law Section: integer (nullable = true)
      -- Sub Division: string (nullable = true)
      -- Days Parking In Effect : string (nullable = true)
      -- From Hours In Effect: string (nullable = true)
      -- To Hours In Effect: string (nullable = true)
      -- Vehicle Color: string (nullable = true)
      -- Unregistered Vehicle?: integer (nullable = true)
      -- Vehicle Year: integer (nullable = true)
      -- Meter Number: string (nullable = true)
      |-- Feet From Curb: integer (nullable = true)
      -- Violation Description: string (nullable = true)
columns_selected = ["Registration State","Plate Type",\
              "Violation Code", "Vehicle Body Type", "Vehicle Make", "Issuing Agency", "Street Code1", \
              "Street Code2", "Street Code3", "Violation Location", "Violation Precinct", \
              "Issuer Precinct", "Issuer Code", "Issuer Command", \
              "Violation County", "Law Section", "Sub Division", "Vehicle Color"]
# removed the Violation_Legal_Code, summons_number , plate_id as it has same values
df = df.select(columns_selected)
# clean up the data as many have incorrect values.
df = df[(df['Registration State'] != "99") & (df['Plate Type'] != "999") \
       & (df['Violation Code'] != 0)]
# Check if the null value still exist
df.select([count(when(isnull(c), c)).alias(c) for c in df.columns]).show()
     |Registration State|Plate Type|Violation Code|Vehicle Body Type|Vehicle Make|Issuing Agency|Street Code1|Street Code2|Street Code3|
    0 0 0 0 4928 3237 0 0 0
                                                                                                                            al
        -----
```

https://colab.research.google.com/drive/1brZk_va1BStF4fLJyzPXKLiYlxxNBRPo#scrollTo=PSpELmcGQRpk&printMode=true

df.toPandas().nunique()

Registration State	60
Plate Type	46
Violation Code	90
Vehicle Body Type	336
Vehicle Make	731
Issuing Agency	12
Street Code1	3422
Street Code2	3719
Street Code3	3648
Violation Location	135
Violation Precinct	136
Issuer Precinct	277
Issuer Code	10952
Issuer Command	602
Violation County	7
Law Section	2
Sub Division	122
Vehicle Color	489
dtvpe: int64	

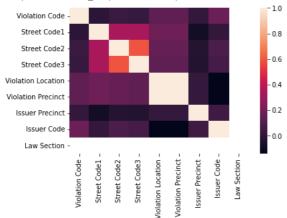
df = df.na.drop()

df.dropDuplicates()

DataFrame[Registration State: string, Plate Type: string, Violation Code: int, Vehicle Body Type: string, Vehicle Make: string, Issuing Agency: string, Street Code1: int, Street Code2: int, Street Code3: int, Violation Location: int, Violation Precinct: int, Issuer Code: int, Issuer Command: string, Violation County: string, Law Section: int, Sub Division: string, Vehicle Color: string]

sns.heatmap(df.toPandas().corr())

<matplotlib.axes._subplots.AxesSubplot at 0x7f4e466aa550>



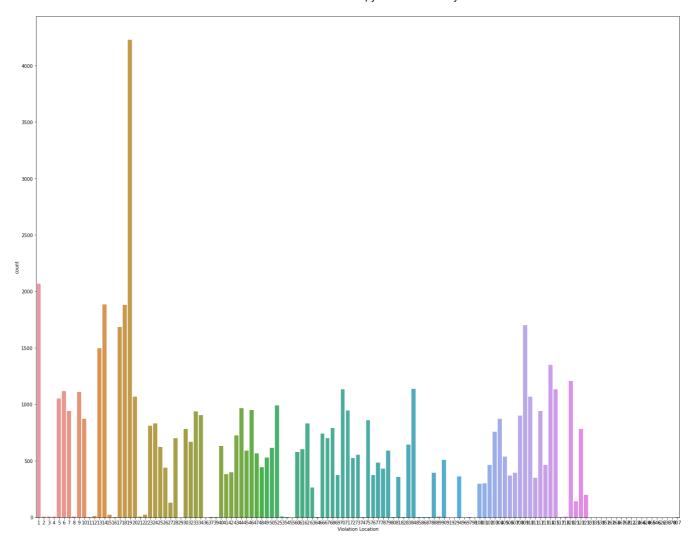
check how labels are ditributed
df.toPandas()['Violation Location'].value_counts()

```
4227
        2070
1
14
        1884
18
        1883
109
        1700
678
92
           1
152
           1
138
           1
```

Name: Violation Location, Length: 127, dtype: int64

df.createOrReplaceTempView("nyc")

```
violation_table = df.groupBy('Violation Location').count()
fig = plt.figure(figsize =(25, 20))
sns.barplot(x='Violation Location',y='count',data=violation_table.toPandas())
plt.show()
```



```
# from pyspark.sql.functions import *
# df=df.withColumn("IssueDate", to_timestamp(col("Issue_Date")))
# type(df)
from pyspark.sql.types import StringType,DoubleType
cols = [F.col(field[0]).cast('double') \ if \ (field[1] == 'int') \ else \ F.col(field[0]) \ for \ field \ in \ df.dtypes]
df = df.select(cols)
df.printSchema()
     root
       |-- Registration State: string (nullable = true)
       -- Plate Type: string (nullable = true)
       |-- Violation Code: double (nullable = true)
       |-- Vehicle Body Type: string (nullable = true)
       -- Vehicle Make: string (nullable = true)
       |-- Issuing Agency: string (nullable = true)
       |-- Street Code1: double (nullable = true)
|-- Street Code2: double (nullable = true)
       |-- Street Code3: double (nullable = true)
       -- Violation Location: double (nullable = true)
       -- Violation Precinct: double (nullable = true)
       -- Issuer Precinct: double (nullable = true)
       |-- Issuer Code: double (nullable = true)
       -- Issuer Command: string (nullable = true)
       -- Violation County: string (nullable = true)
       |-- Law Section: double (nullable = true)
       -- Sub Division: string (nullable = true)
       |-- Vehicle Color: string (nullable = true)
```

```
label_count=len(df.select('Violation Location').distinct().collect())
print(label_count)
```

Convert categorical data to numerical then to vector assemle the data for ML lib

```
# Split the data into categorical and continuos to use it pipeline
cat_columns = [x for (x, dataType) in df.dtypes if dataType == 'string']
num_columns = [x for (x, dataType) in df.dtypes if ((dataType != 'string') & (x != 'Violation_Location'))]
# convert categorical to string indexer then to one hot encoder using pipeline
from pyspark.ml.feature import StringIndexer
from pyspark.ml.feature import OneHotEncoder
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.feature import MinMaxScaler
from pyspark.ml import Pipeline
from pyspark.ml import PipelineModel
#1. Convert to string indexer
indexers = [ StringIndexer(inputCol=c, handleInvalid="skip",outputCol="{0} indexed".format(c)) for c in cat columns ]
#2. Convert to OnehotEncoder
encoders = [ OneHotEncoder(
                inputCol=indexer.getOutputCol(),
                outputCol="{0}_encoded".format(indexer.getOutputCol()))
             for indexer in indexers ]
#3. Assembler
assembler = VectorAssembler(inputCols=[encoder.getOutputCol() for encoder in encoders]
                                + num_columns, outputCol="features_vect")
## MinMax Scalar
scaler = MinMaxScaler(inputCol="features_vect", outputCol="features_scaled")
pipeline = Pipeline(stages=indexers + encoders+ [assembler]+[scaler])
pipeline_model = pipeline.fit(df)
# save the model for fitting the data for prediction
pipeline_model.save("va_model")
df = pipeline_model.transform(df)
Note that above data transformation pipeline is saved as model to transform the new data for prediction
# There are more labels than standard multi label classifier support to predict. So lets index the label.
label_indexer = StringIndexer(inputCol="Violation Location", outputCol="VL_indexed")
label_indexers_model = label_indexer.fit(df)
df= label_indexers_model.transform(df)
                                                         + Code — + Text
# Display the features column to check whether its dense or sparse vector
df.select("features_scaled","VL_indexed").show(3)
          features_scaled|VL_indexed|
     |(1926,[0,60,99,30...|
                                23.0
     [(1926, [0,60,99,30...]
                                23.0
     |(1926,[0,60,104,3...|
                                23.0
     only showing top 3 rows
# from pyspark.ml.linalg import Vectors, VectorUDT
# # define function to convert parse to dense vector
# sparseToDense = F.udf(lambda v : Vectors.dense(v), VectorUDT())
# densefeatureDF = df.withColumn('dense_features', sparseToDense('features_scaled'))
# densefeatureDF.select("features_scaled", "dense_features").show(1)
data = df.select(F.col("features_scaled").alias("features"),
                          F.col("VL indexed").alias("label"))
```

```
# cache and persist
from pyspark import StorageLevel
# persist the dataframe in both memo
persist_data = data.persist(StorageLevel.MEMORY_AND_DISK)
trainDF, testDF = persist_data.randomSplit([0.7,0.3], seed = 2642)
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
# rf_model = RandomForestClassifier(featuresCol='features', labelCol='label',
                     predictionCol='prediction',maxDepth=30,
                      impurity='gini', subsamplingRate= .5).fit(trainDF)
# pred = rf_model.transform(testDF)
# pred.show(3)
rf = RandomForestClassifier(featuresCol='features', labelCol='label',
                    predictionCol='prediction', impurity='gini')
paramGrid = ParamGridBuilder()\
    .addGrid(rf.maxDepth, [30]) \
    .addGrid(rf.subsamplingRate, [0.1, 0.5])\
    .addGrid(rf.numTrees, [20,40])\
    .build()
evaluator=MulticlassClassificationEvaluator(predictionCol="prediction")
rf tvs = TrainValidationSplit(estimator=rf,
                           estimatorParamMaps=paramGrid,
                           evaluator=evaluator,
                           # 80% of the data will be used for training, 20% for validation.
rf tvs model = rf tvs.fit(trainDF)
rf_predictonDF = rf_tvs_model.transform(testDF)
acc = evaluator.evaluate(rf_predictonDF)
print("Prediction Accuracy: ", acc)
rf_predictonDF.select("features","label", "prediction").show(5)
     Prediction Accuracy: 0.8608196539256541
                features|label|prediction|
     |(1926,[0,59,98,30...| 0.0|
     |(1926,[0,59,98,30...| 24.0|
                                       24.0
     |(1926,[0,59,98,30...| 34.0|
                                       40.0
     |(1926,[0,59,98,30...| 23.0|
                                       24.0
     |(1926,[0,59,98,30...| 24.0|
                                       23.01
     only showing top 5 rows
from pyspark.ml.feature import IndexToString
from pyspark.sql.types import DoubleType
from pyspark.ml import PipelineModel
import os
dirname = os.getcwd()
print('path=',dirname)
# get the best model and save it
lr_model = rf_tvs_model.bestModel
filename = "rf_model"
model_path= os.path.join(dirname, filename)
print(model path)
# Index to String converter to get the location
label_conv = IndexToString(inputCol="prediction", outputCol="Violation_Location", labels=label_indexers_model.labels)
# save the Randomforest regression model as pipeline model
pm = PipelineModel(stages=[lr_model,label_conv])
# Write the model
pm.save(filename)
     path= /content
     /content/rf_model
```

^{*}Best model from TrainValidationSplit is stored along with converion for prediction to violation location using index to string. *

```
from google.colab import files
!zip -r 'rf_model.zip' 'rf_model'
!zip -r 'va_model.zip' 'va_model'
files.download('/content/rf_model.zip')
files.download('/content/va_model.zip')
from sklearn.metrics import confusion_matrix
from pyspark.ml.feature import IndexToString
from pyspark.sql.types import DoubleType
label_conv = IndexToString(inputCol="prediction", outputCol="Violation_Location", labels=label_indexers_model.labels)
pred_label = label_conv.transform(rf_predictonDF)
# convert Violation_Location to DoubleType for accuracy calulation.
pred_label = pred_label.withColumn("Violation_Location", col("Violation_Location").cast(DoubleType()))
evaluator = \texttt{MulticlassClassificationEvaluator} (\texttt{predictionCol="prediction"})
acc = evaluator.evaluate(pred_label)
print("Prediction Accuracy: ", acc)
# import PipelineModel from pyspark.ml package
from pyspark.ml import PipelineModel
import os
import zipfile
dirname = os.getcwd()
print('path=',dirname)
rf_filename = "rf_model"
model_path= os.path.join(dirname, rf_filename)
print(model_path)
# load the model from the location it is stored
# The loaded model acts as PipelineModel
pipeline = PipelineModel.load(model path)
# use the PipelineModel object to perform prediciton on test data.
\ensuremath{\mbox{\sc #}} Use .transform() to perfrom prediction
prediction = pipeline.transform(testDF)
# print the results
# prediction.select('label', 'prediction', 'Violation_Location').show(5)
prediction.select('prediction','Violation_Location').show(5)
spark.stop()
```

Fit the model and check the accuracy

• x